

DEEP LEARNING COM TENSORFLOW



# SEQUÊNCIAS

DIEGO RODRIGUES DSC

INFNET

# CRONOGRAMA

<b>Dia</b>	<b>Aula</b>	<b>Trab</b>
<del>02/09</del>	<del>Workshop de Deep Learning</del>	
<del>04/09</del>	<del>Deep FeedForward</del>	
<del>09/09</del>	<del>Rede Neural Convolutiva</del>	<del>Modelo Baseline</del>
<del>11/09</del>	<del>AutoEncoders</del>	
16/09	<del>Representation &amp; Transfer Learning</del>	<del>Modelo Profundo</del>
18/09	Sequências	
23/09	Modelos Generativos	Deployment
25/09	Apresentação dos Trabalhos Parte II	



# SEQUÊNCIAS

- PARTE 1 : TEORIA

- BUSINESS UNDERSTANDING

- ESTADO

- FILTRO DE KALMAN

- MODELING

- REDES NEURAIS RECORRENTES

- PROBLEMA DO GRADIENTE

- PARTE 2 : PRÁTICA

- SEQUENCIAS

- PARTE 3 : TRABALHOS

The background is a light gray gradient. In the top-left and bottom-right corners, there are several realistic water droplets of various sizes, rendered with highlights and shadows to give them a 3D appearance. A faint, circular watermark is visible in the center of the slide, containing a play button icon and the text '© 2014 KAPPA'.

# PARTE 1 : TEORIA

The image features a light gray background with a subtle radial gradient. In the top-left and bottom-right corners, there are clusters of realistic, three-dimensional water droplets of various sizes. A faint, circular watermark is visible in the upper center of the page.

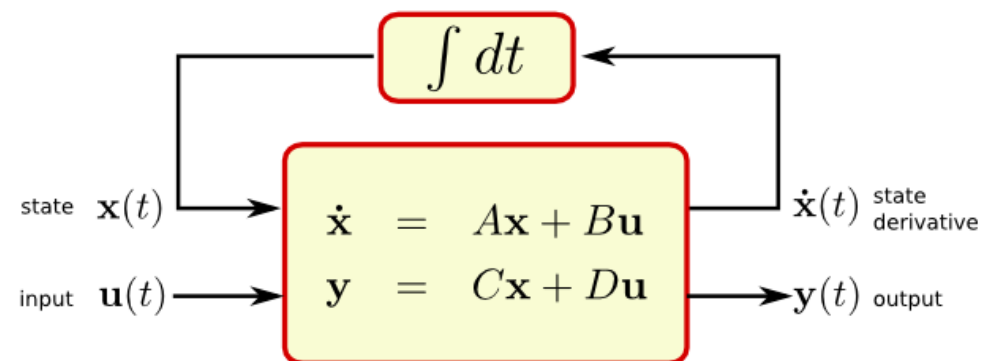
# **BUSINESS UNDERSTANDING**

# ESTADO

O conceito de estado se refere a um **conjunto de variáveis** que **capturam todas as informações necessárias** para **descrever completamente o status atual de um sistema**. O estado contém **todas as informações sobre o passado** do sistema que são **relevantes para prever seu comportamento futuro**.

Em termos práticos, **controlar um sistema envolve manipular as entradas para direcioná-lo a um estado desejado**, respondendo adequadamente **às perturbações e mudanças no ambiente**. **Entender e modelar o estado** é fundamental para projetar controladores eficientes e para a estabilidade e desempenho do sistema.


state space



System represented as a collection of coupled linear first-order differential equations.

coupled linear first-order differential equations.  
system represented as a collection of

# ESPAÇO DE ESTADOS



## System type

## State-space model

Continuous time-invariant

$$\dot{\mathbf{x}}(t) = \mathbf{A}\mathbf{x}(t) + \mathbf{B}\mathbf{u}(t)$$

$$\mathbf{y}(t) = \mathbf{C}\mathbf{x}(t) + \mathbf{D}\mathbf{u}(t)$$

Continuous time-variant

$$\dot{\mathbf{x}}(t) = \mathbf{A}(t)\mathbf{x}(t) + \mathbf{B}(t)\mathbf{u}(t)$$

$$\mathbf{y}(t) = \mathbf{C}(t)\mathbf{x}(t) + \mathbf{D}(t)\mathbf{u}(t)$$

Explicit discrete time-invariant

$$\mathbf{x}(k+1) = \mathbf{A}\mathbf{x}(k) + \mathbf{B}\mathbf{u}(k)$$

$$\mathbf{y}(k) = \mathbf{C}\mathbf{x}(k) + \mathbf{D}\mathbf{u}(k)$$

Explicit discrete time-variant

$$\mathbf{x}(k+1) = \mathbf{A}(k)\mathbf{x}(k) + \mathbf{B}(k)\mathbf{u}(k)$$

$$\mathbf{y}(k) = \mathbf{C}(k)\mathbf{x}(k) + \mathbf{D}(k)\mathbf{u}(k)$$

Laplace domain of  
continuous time-invariant

$$s\mathbf{X}(s) - \mathbf{x}(0) = \mathbf{A}\mathbf{X}(s) + \mathbf{B}\mathbf{U}(s)$$

$$\mathbf{Y}(s) = \mathbf{C}\mathbf{X}(s) + \mathbf{D}\mathbf{U}(s)$$

Z-domain of  
discrete time-invariant

$$z\mathbf{X}(z) - z\mathbf{x}(0) = \mathbf{A}\mathbf{X}(z) + \mathbf{B}\mathbf{U}(z)$$

$$\mathbf{Y}(z) = \mathbf{C}\mathbf{X}(z) + \mathbf{D}\mathbf{U}(z)$$

discrete time-invariant

$$\mathbf{X}(z) = \mathbf{C}\mathbf{X}(z) + \mathbf{D}\mathbf{U}(z)$$

Z-domain of

$$z\mathbf{X}(z) - z\mathbf{x}(0) = \mathbf{A}\mathbf{X}(z) + \mathbf{B}\mathbf{U}(z)$$

continuous time-invariant


$$\mathbf{X}(s) = \mathbf{C}\mathbf{X}(s) + \mathbf{D}\mathbf{U}(s)$$

continuous time-variant


$$\dot{\mathbf{x}}(t) = \mathbf{A}(t)\mathbf{x}(t) + \mathbf{B}(t)\mathbf{u}(t)$$



# SISTEMA 1



Freebody Diagram



Equation

$$m \cdot \ddot{x} + k_1 x + k_2 x - k_1 z = f_a$$

$$b \cdot \dot{z} + k_1 z - k_1 x = 0$$

$$\begin{aligned}\dot{q}_1 &= \dot{x} = q_2 \\ \dot{q}_2 &= \ddot{x} = \frac{1}{m}(f_a - k_1 x - k_2 x + k_1 z) \\ &= \frac{1}{m}(f_a - k_1 q_1 - k_2 q_1 + k_1 q_3) \\ \dot{q}_3 &= \dot{z} = \frac{k_1}{b}(x - z) = \frac{k_1}{b}(q_1 - q_3)\end{aligned}$$


$$\begin{aligned}\dot{\mathbf{q}} &= \mathbf{A}\mathbf{q} + \mathbf{B}u \\ y &= \mathbf{C}\mathbf{q} + \mathbf{D}u\end{aligned}\quad \mathbf{A} = \begin{bmatrix} 0 & 1 & 0 \\ -\frac{k_1 + k_2}{m} & 0 & \frac{k_1}{m} \\ \frac{k_1}{b} & 0 & -\frac{k_1}{b} \end{bmatrix} \quad \mathbf{B} = \begin{bmatrix} 0 \\ \frac{1}{m} \\ 0 \end{bmatrix}$$

$$\mathbf{C} = [0 \quad 0 \quad 1] \quad \mathbf{D} = 0$$



# SISTEMA 2

## State-Space Example 1



### system dynamics

$$m\ddot{x} + b\dot{x} + kx = f$$

$$\text{input } u = f$$

$$\text{output } y = f - b\dot{x} - kx$$

### transfer function

$$\frac{Y(s)}{U(s)} = \frac{ms^2}{ms^2 + bs + k}$$

### state-space modeling

$$\text{state variables: } x_1 = x, x_2 = \dot{x}$$

$x_1, x_2$ : second-order system  
can completely describe  
the system's state

### state-space dynamics

$$\begin{bmatrix} \dot{x}_1 \\ \dot{x}_2 \end{bmatrix} = \begin{bmatrix} 0 & 1 \\ -k/m & -b/m \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} + \begin{bmatrix} 0 \\ 1/m \end{bmatrix} u$$

first-order vector  
differential equation

$$y = [-k \quad -b] \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} + 1 \times u$$

### state-space representation

$$\dot{x} = Ax + Bu, \quad y = Cx + Du$$

# SISTEMA 3

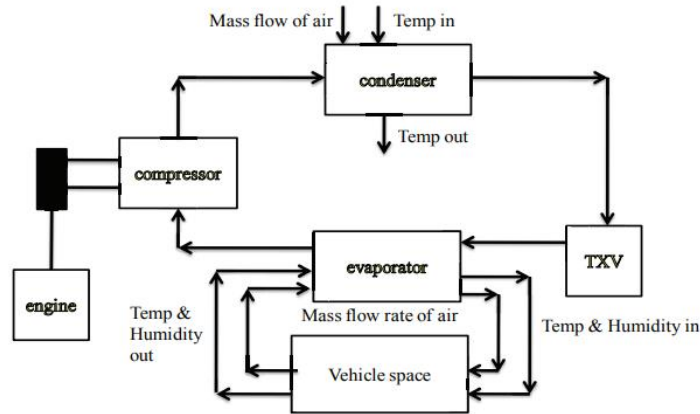


Figure 1. Automotive Air Conditioning System Block Diagram Representation




Figure 2. Physical Model of Passengers' Compartment and Thermal Loads

$$h_1(x,u) = \begin{bmatrix} -C_{pe} \rho f (T_m - T_s) + U_0 A_0 (T_a - T_r) + C_{pa} \rho f (T_a - T_r) + K_{spl} f \\ -h_{im} A_{im} (T_{im} - T_r) \\ C_p \rho f (T_d - T_s) + \alpha_1 A_1 \left( T_w - \frac{T_m + T_d}{2} \right) \\ C_p \rho f (T_d - T_s) + \rho f h_{fg} (W_m - W_s) + \alpha_2 A_2 \left( T_w - \frac{T_d + T_s}{2} \right) \\ \alpha_1 A_1 \left( \frac{T_m + T_d}{2} - T_w \right) + \alpha_2 A_2 \left( \frac{T_d + T_s}{2} - T_w \right) - M_{ref} (h_{r2} - h_{r1}) \\ -\rho f h_{fg} (W_m - W_s) + M_r h_{fg} (W_a - W_r) + Q_{pl} \\ 0 \end{bmatrix} \quad - (14.a), \quad h_2(z) = \begin{bmatrix} Q_s + Q_{ps} \\ 0 \\ 0 \\ 0 \\ 0 \\ Q_{pl} \\ 0 \end{bmatrix} \quad - (14.b)$$

$$D = \begin{bmatrix} M_r C_{pr} & M_{im} C_c & 0 & 0 & 0 & 0 & 0 \\ 0 & M_{im} C_c & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & C_p \rho V_{h1} & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & C_p \rho V_{h2} & 0 & 0 & \rho V_{h2} h_{fg} \\ 0 & 0 & 0 & 0 & (C_p \rho V)_w & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & M_r h_{fg} & 0 \\ 0 & 0 & 0 & \left( \frac{-2 \times 0.0198 T_s + 0.085}{1000} \right) & 0 & 0 & 1 \end{bmatrix} \quad - (14.c)$$

# FILTRO DE KALMAN

O **Filtro de Kalman** é um algoritmo recursivo **que fornece estimativas de estados ocultos** de um sistema dinâmico com base em medições ruidosas ao longo do tempo.

Utilizado amplamente em controle de sistemas, navegação, robótica, e **previsão de séries temporais**.

**Cenário:** Desejamos **rastrear o estado de um sistema** (como a posição e velocidade de um objeto) que está sujeito a:

- Dinâmica do sistema (movimento).
- Ruído nos sensores que medem essas variáveis.
- **Exemplo:** Estimar a posição de um veículo usando um sensor GPS ruidoso.

The diagram illustrates the two fundamental equations of a Kalman filter, with annotations identifying the mathematical types of the variables:

$$\mathbf{x}'(t) = \mathbf{A}\mathbf{x}(t) + \mathbf{B}u(t)$$
$$y(t) = \mathbf{C}\mathbf{x}(t) + \mathbf{D}u(t)$$

Annotations for the equations:

- $\mathbf{x}'(t)$ : Vector
- $\mathbf{A}$ : Matrix
- $\mathbf{x}(t)$ : Vector
- $\mathbf{B}$ : Vector or Matrix
- $u(t)$ : Scalar
- $y(t)$ : Scalar
- $\mathbf{C}$ : Vector or Matrix
- $\mathbf{x}(t)$ : Vector
- $\mathbf{D}$ : Vector or Matrix
- $u(t)$ : Scalar

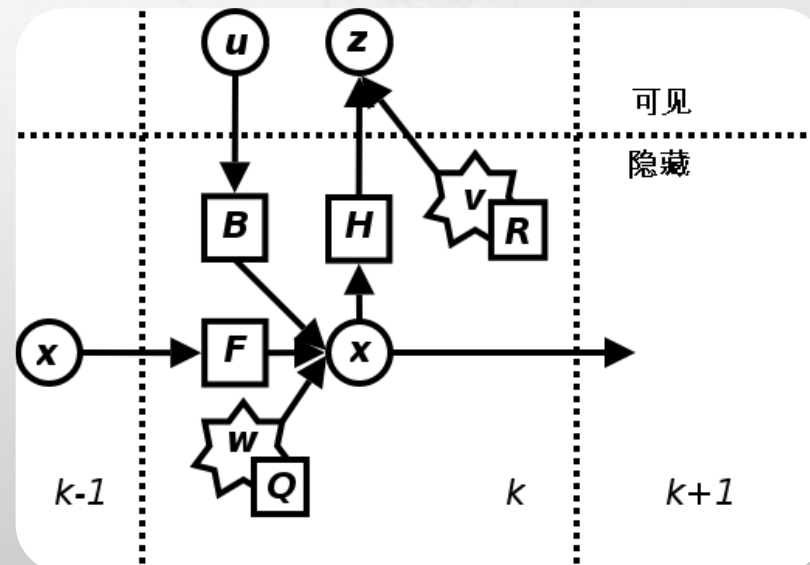

# FILTRO DE KALMAN

## Predição

- Estima o próximo estado com base no estado atual e no modelo de transição.
- Atualiza a incerteza associada à predição.

## Atualização

- Com base em novas observações, corrige a predição do estado.
- Calcula o ganho de Kalman (peso que indica a confiança na observação vs. a predição).
- Atualiza o estado estimado e reduz a incerteza.



# MODELING




# APLICAÇÕES

**Processamento de Linguagem Natural (NLP):** Modelagem de dependências entre palavras em uma frase.

**Previsão de Séries Temporais:** Como em previsão de demanda de produtos, mercado financeiro, etc.

**Reconhecimento de Fala:** Mapeamento de uma sequência de áudio para texto.



# SÉRIES TEMPORAIS

$$y_t = T_t + C_t + S_t + I_t,$$

$$y_t = T_t \times C_t \times S_t \times I_t.$$

$T_t$  : média / média móvel

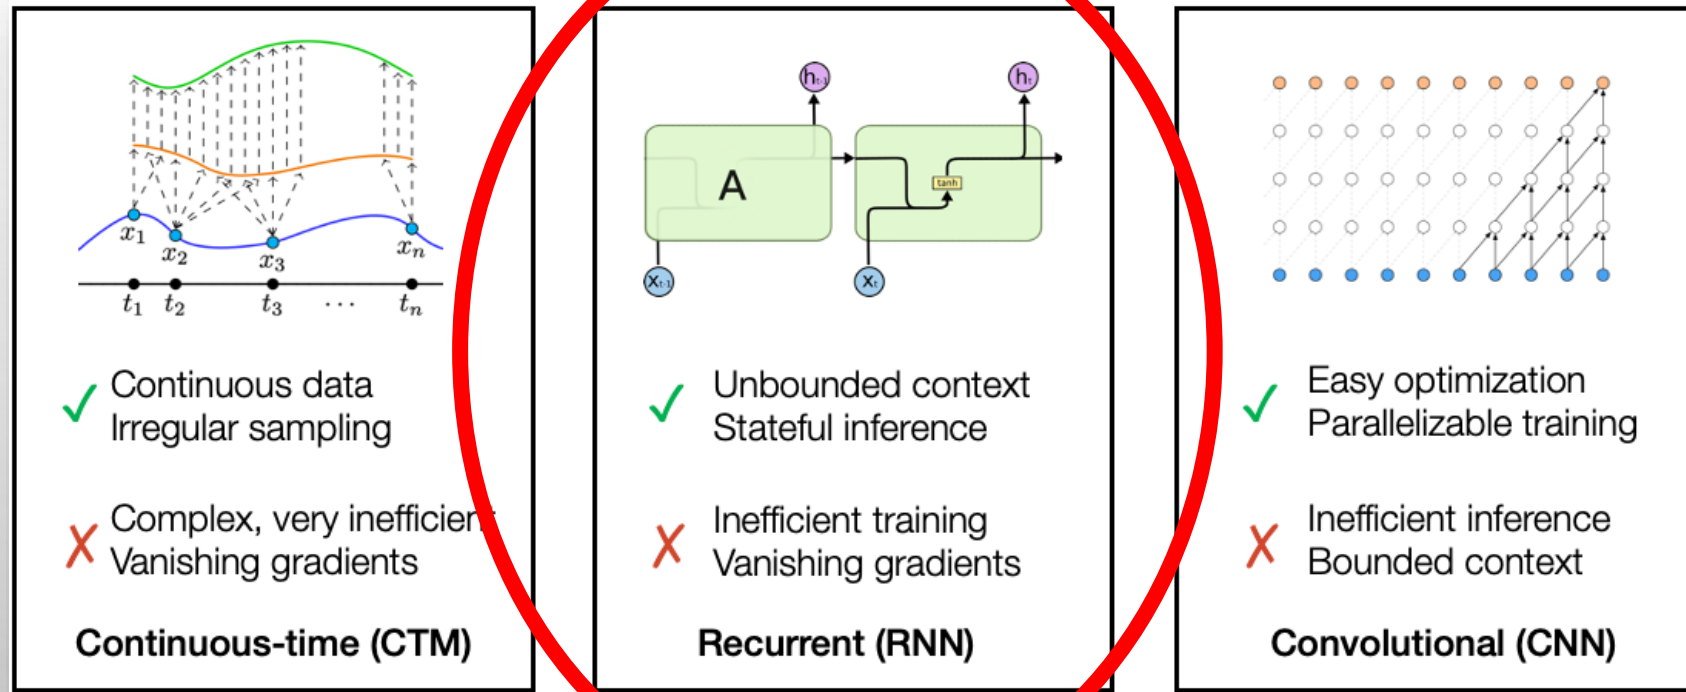
$C_t$ : Fourier Passa Baixas / Média Móvel

$S_t$ : Fourier Passa Altas / Picos Remanescentes

$I_t$ : auto-regressão / rede neural



# SÉRIES TEMPORAIS COM REDES NEURAIS




# ESPAÇO DE REPRESENTAÇÕES


Largura

Complexidade vs Figura de Mérito

Artificial Neural Network



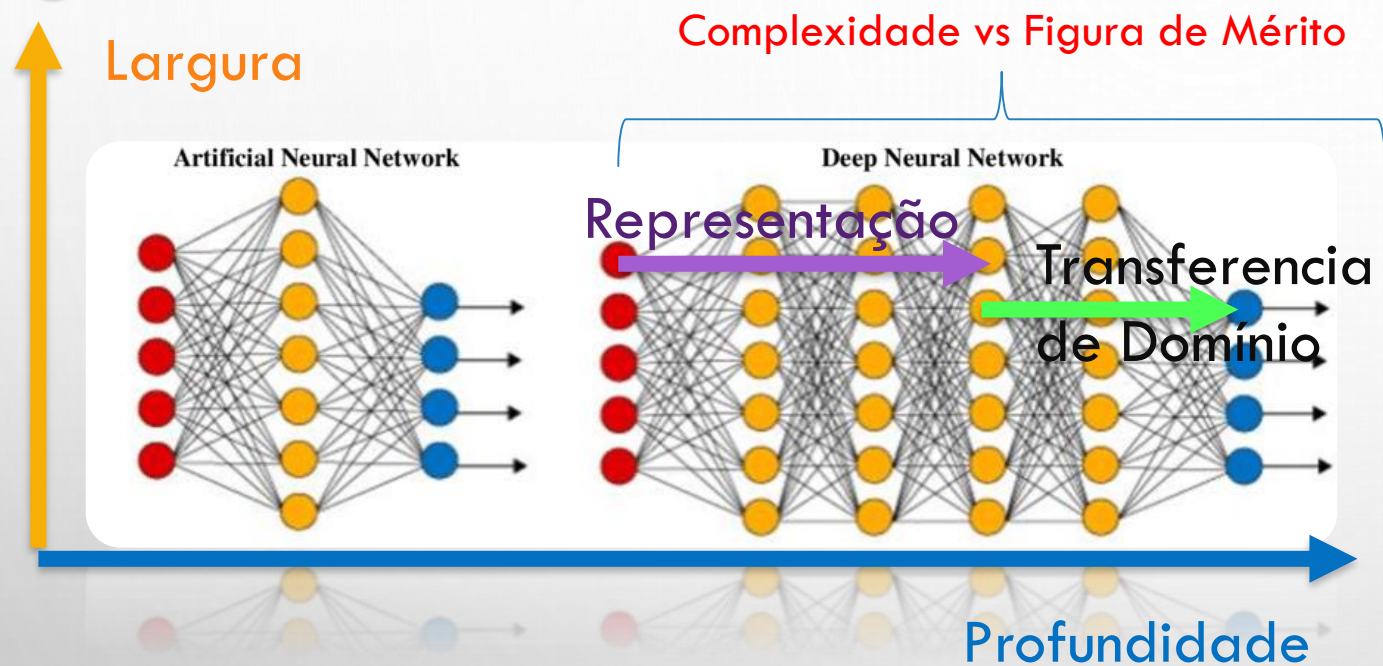
Deep Neural Network



Profundidade




# VERSÃO 1, VERSÃO 2, ETC




# TIPOS DE JANELAMENTO


one to one




one to many




many to one



many to many



many to many





# REDE NEURAL RECORRENTE

**Redes Neurais Recorrentes** são um tipo de rede neural projetada para lidar com **dados sequenciais**.

A **principal característica** de uma RNN é a **retroalimentação**: a saída de uma unidade de tempo influencia a entrada da próxima unidade de tempo.

Isso permite que a RNN **mantenha um estado interno** que **armazena informações sobre entradas anteriores**, tornando-a adequada para tarefas como **séries temporais**, **processamento de linguagem natural (NLP)**, e **reconhecimento de fala**.

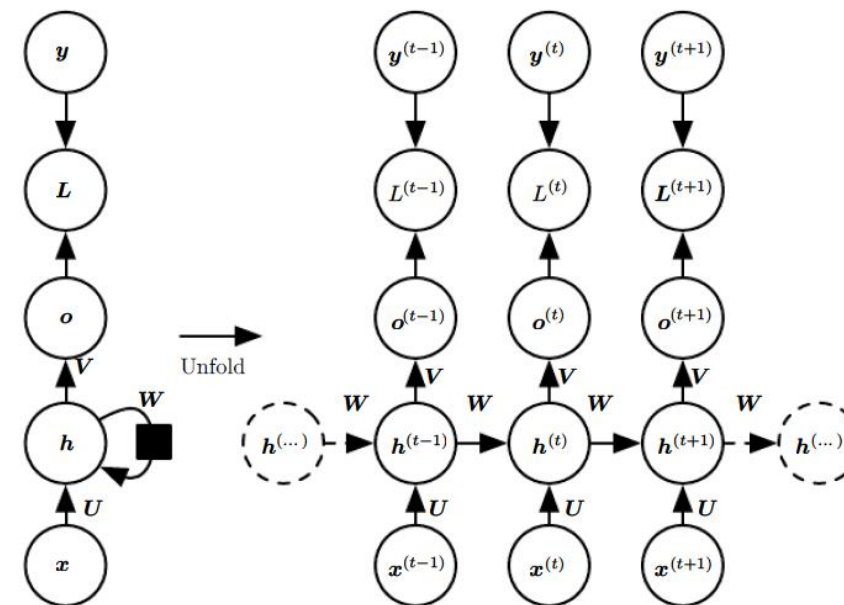


Figure 10.3: The computational graph to compute the training loss of a recurrent network that maps an input sequence of  $x$  values to a corresponding sequence of output  $o$  values. A loss  $L$  measures how far each  $o$  is from the corresponding training target  $y$ . When using softmax outputs, we assume  $o$  is the unnormalized log probabilities. The loss  $L$  internally computes  $\hat{y} = \text{softmax}(o)$  and compares this to the target  $y$ . The RNN has input to hidden connections parametrized by a weight matrix  $U$ , hidden-to-hidden recurrent connections parametrized by a weight matrix  $W$ , and hidden-to-output connections parametrized by a weight matrix  $V$ . Equation 10.8 defines forward propagation in this model. (Left) The RNN and its loss drawn with recurrent connections. (Right) The same seen as an time-unfolded computational graph, where each node is now associated with one particular time instance.

## REDE NEURAL RECORRENTE II




Figure 10.4: An RNN whose only recurrence is the feedback connection from the output to the hidden layer. At each time step  $t$ , the input is  $x_t$ , the hidden layer activations are  $h^{(t)}$ , the outputs are  $o^{(t)}$ , the targets are  $y^{(t)}$  and the loss is  $L^{(t)}$ . (Left) Circuit diagram. (Right) Unfolded computational graph. Such an RNN is less powerful (can express a smaller set of functions) than those in the family represented by figure 10.3. The RNN in figure 10.3 can choose to put any information it wants about the past into its hidden representation  $h$  and transmit  $h$  to the future. The RNN in this figure is trained to put a specific output value into  $o$ , and  $o$  is the only information it is allowed to send to the future. There are no direct connections from  $h$  going forward. The previous  $h$  is connected to the present only indirectly, via the predictions it was used to produce. Unless  $o$  is very high-dimensional and rich, it will usually lack important information from the past. This makes the RNN in this figure less powerful, but it may be easier to train because each time step can be trained in isolation from the others, allowing greater parallelization during training, as described in section 10.2.1.

## REDE NEURAL RECORRENTE III

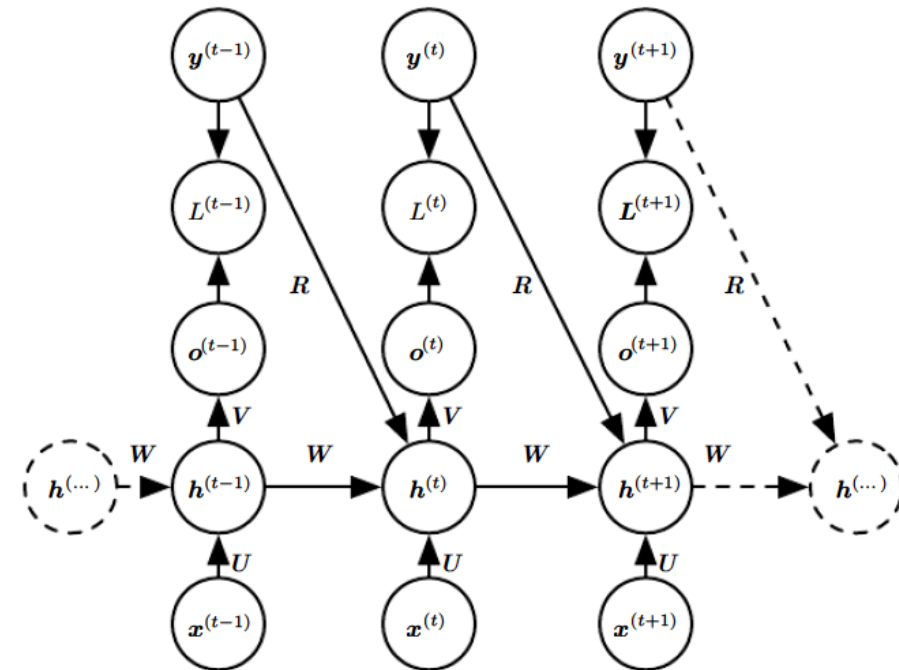


Figure 10.10: A conditional recurrent neural network mapping a variable-length sequence of  $x$  values into a distribution over sequences of  $y$  values of the same length. Compared to figure 10.3, this RNN contains connections from the previous output to the current state. These connections allow this RNN to model an arbitrary distribution over sequences of  $y$  given sequences of  $x$  of the same length. The RNN of figure 10.3 is only able to represent distributions in which the  $y$  values are conditionally independent from each other given the  $x$  values.



# REDE NEURAL RECORRENTE BIDIRECIONAL




Figure 10.11: Computation of a typical bidirectional recurrent neural network, meant to learn to map input sequences  $\mathbf{x}$  to target sequences  $\mathbf{y}$ , with loss  $L^{(t)}$  at each step  $t$ . The  $\mathbf{h}$  recurrence propagates information forward in time (towards the right) while the  $\mathbf{g}$  recurrence propagates information backward in time (towards the left). Thus at each point  $t$ , the output units  $o^{(t)}$  can benefit from a relevant summary of the past in its  $h^{(t)}$  input and from a relevant summary of the future in its  $g^{(t)}$  input.

# REDE NEURAL ENCODER RECORRENTE DECODER




Figure 10.12: Example of an encoder-decoder or sequence-to-sequence RNN architecture, for learning to generate an output sequence  $(y^{(1)}, \dots, y^{(n_y)})$  given an input sequence  $(x^{(1)}, x^{(2)}, \dots, x^{(n_x)})$ . It is composed of an encoder RNN that reads the input sequence and a decoder RNN that generates the output sequence (or computes the probability of a given output sequence). The final hidden state of the encoder RNN is used to compute a generally fixed-size context variable  $C$  which represents a semantic summary of the input sequence and is given as input to the decoder RNN.

# REDE RECORRENTE PROFUNDA




Figure 10.13: A recurrent neural network can be made deep in many ways (Pascanu *et al.*, 2014a). (a) The hidden recurrent state can be broken down into groups organized hierarchically. (b) Deeper computation (e.g., an MLP) can be introduced in the input-to-hidden, hidden-to-hidden and hidden-to-output parts. This may lengthen the shortest path linking different time steps. (c) The path-lengthening effect can be mitigated by introducing skip connections.

# REDE RECORRENTE RECURSIVA




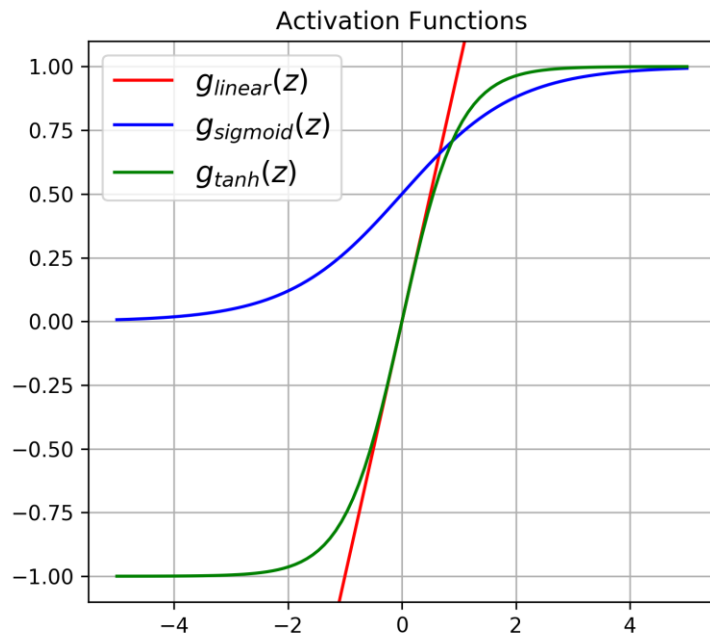


Figure 10.14: A recursive network has a computational graph that generalizes that of the recurrent network from a chain to a tree. A variable-size sequence  $\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(t)}$  can be mapped to a fixed-size representation (the output  $\mathbf{o}$ ), with a fixed set of parameters (the weight matrices  $\mathbf{U}$ ,  $\mathbf{V}$ ,  $\mathbf{W}$ ). The figure illustrates a supervised learning case in which some target  $\mathbf{y}$  is provided which is associated with the whole sequence.



# O PROBLEMA DA DISSIPAÇÃO DO GRADIENTE

Some Common Activation Functions & Their Derivatives




# GRADIENT CLIPPING

**O que é:** Técnica para limitar a magnitude dos gradientes durante o treinamento de redes neurais.

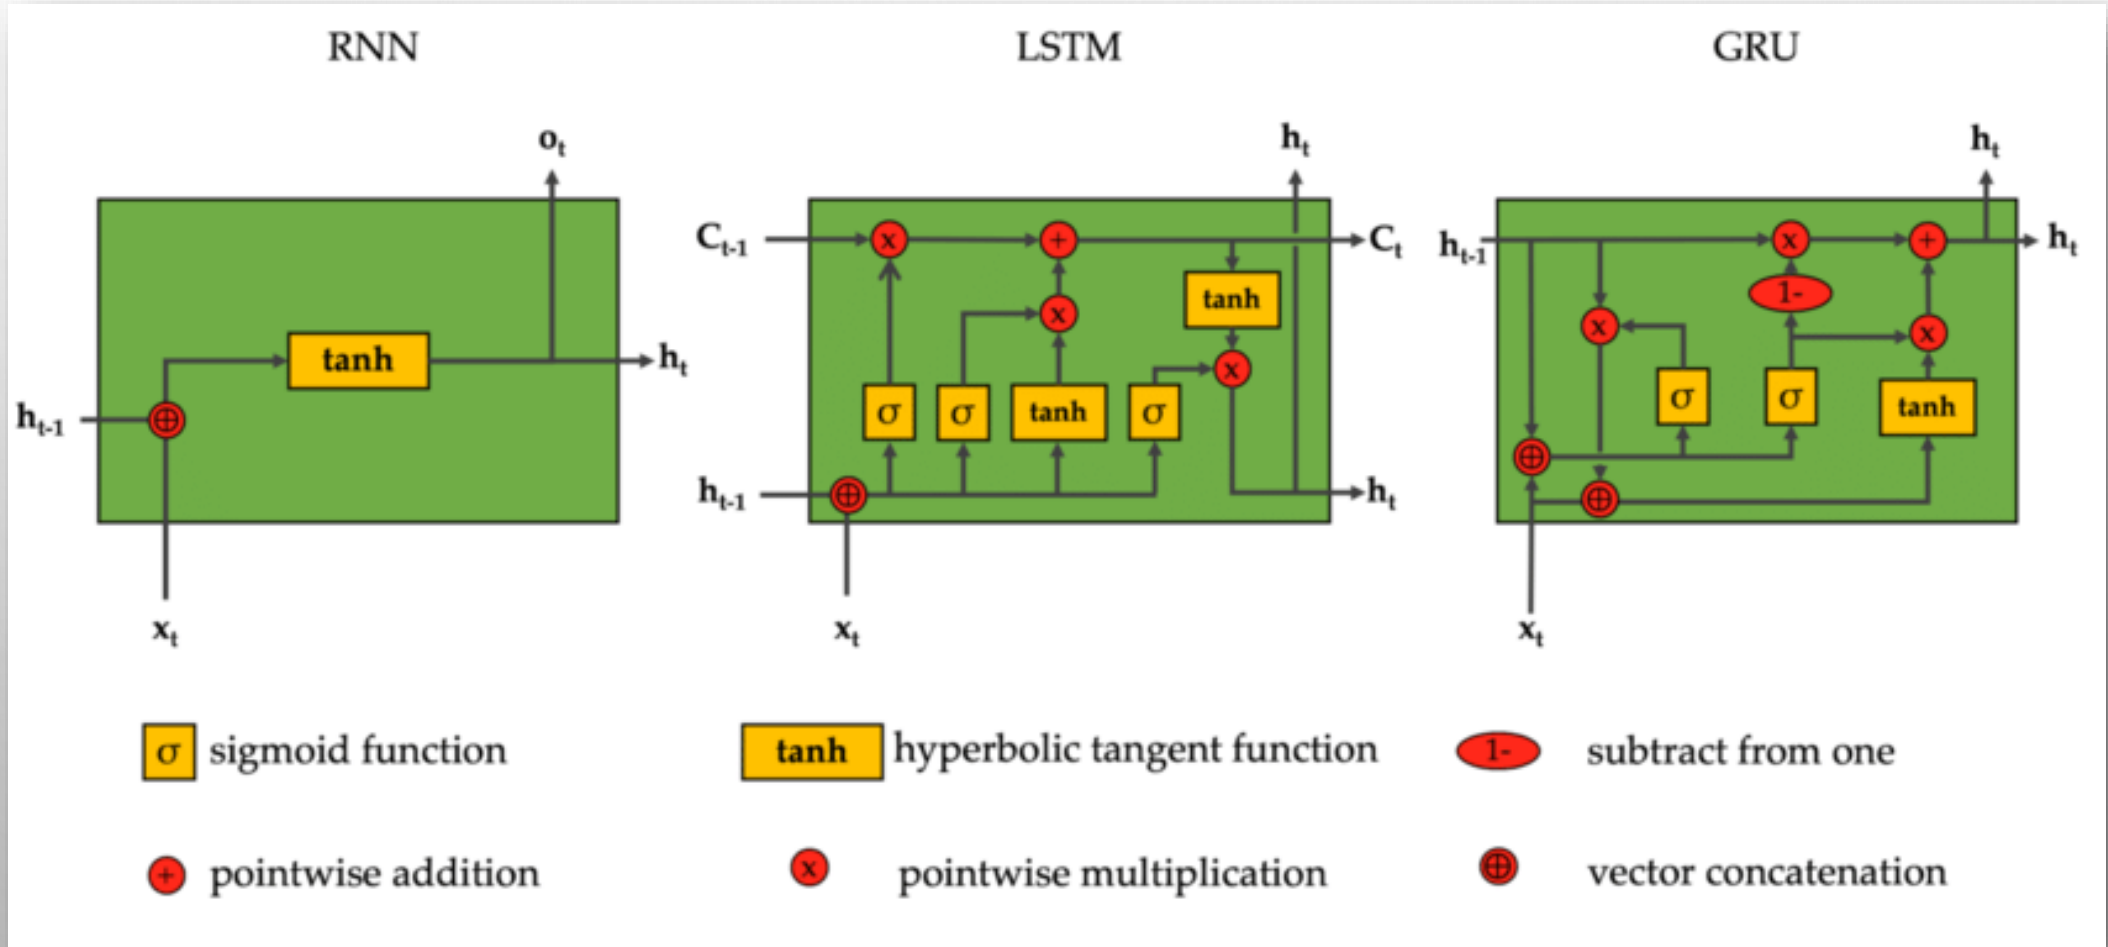
**Por que usar:** Evita a explosão do gradiente em RNNs, melhorando a estabilidade do treinamento.

**Como funciona:**

- Clipping pela Norma (clipnorm): Limita a norma L2 do gradiente a um valor máximo.
- Clipping pelo Valor (clipvalue): Limita o valor absoluto de cada componente do gradiente.



# CONTROLE DO GRADIENTE





# LSTM





Figure 10.16: Block diagram of the LSTM recurrent network “cell.” Cells are connected recurrently to each other, replacing the usual hidden units of ordinary recurrent networks. An input feature is computed with a regular artificial neuron unit. Its value can be accumulated into the state if the sigmoidal input gate allows it. The state unit has a linear self-loop whose weight is controlled by the forget gate. The output of the cell can be shut off by the output gate. All the gating units have a sigmoid nonlinearity, while the input unit can have any squashing nonlinearity. The state unit can also be used as an extra input to the gating units. The black square indicates a delay of a single time step.

# GRU



The background is a light gray gradient. In the top-left and bottom-right corners, there are several realistic water droplets of various sizes, some overlapping. A faint, circular watermark is visible in the upper center of the page.

## PARTE 2 : PRÁTICA

# AMBIENTE PYTHON



6. Machine Learning




6. Deployment



4. Variáveis Aleatórias

5. Visualização



1. Editor de Código



2. Gestor de Ambiente



3. Ambiente Python do Projeto



3. Notebook Dinâmico

# WORKSHOP

- ~~• QUAL A TOPOLOGIA DE DEEP LEARNING ADEQUADA PARA O MEU TRABALHO?~~
- ~~• QUAL CAPÍTULO DO LIVRO MELHOR SE ENQUADRA NO MEU TRABALHO?~~
- ~~• AULA 3: NOVO CICLO DE BUSINESS UNDERSTANDING / GRUPO + MODELO BASELINE TREINADO~~
- **AULA 5 OU 7: MODELO PROFUNDO TREINADO**
- **AULA 7: DEPLOYMENT DO MODELO\***
- **AULA 3-7 > APRESENTAÇÃO TEÓRICA DA(S) TOPOLOGIA(S) + LEITURA DE ARTIGO + ACOMPANHAMENTO DOS TRABALHOS + DEEP DIVE NO CÓDIGO (POR GRUPO)**
- **APRESENTAÇÃO FINAL DOS TRABALHOS**

*\*Passo Opcional*